



Application Programming Interfaces (APIs)

Adopting the NIAID Blueprint

SDSC
SAN DIEGO SUPERCOMPUTER CENTER

GO FAIR US **GO FAIR**
oundation

NCAR | NATIONAL CENTER FOR
ATMOSPHERIC RESEARCH

Knowledge Motifs LLC
Mapping sensible data relationships

Hosted By



[Office of Data Science and Emerging Technologies \(ODSET\)](#)

ODSET's mission is to collaborate strategically across divisions to advance the NIAID mission through data infrastructure and data innovation and data-driven solutions.



Based at the San Diego Supercomputer Center (SDSC), [GO FAIR US \(GFU\)](#) fosters a collaborative community where FAIR approaches can be shared and advanced. GFU promotes practical implementation of FAIR across research domains.


[NIAID Data Landscaping and FAIRification Project](#)

Since March 2025, ODSET and GFU have partnered with NIAID-funded repositories to identify opportunities for increasing data discovery and reuse through Blueprint implementation. This collaboration bridges technical standards with real-world research needs.

ODSET & GFU work with NIAID repositories to enhance data sharing, reuse, and reproducibility

Goals:

- 1) lower barriers to data access and reuse
- 2) provide feedback, webinars, and community of practice to (try to) connect NIAID data



A Blueprint for Including Digital Objects in the NIAID Data Ecosystem:

1. Minimal Metadata Schema
2. Persistent Identifiers
3. Minimal API Specification
4. Citation Guidance



Read the NIAID Blueprint
(DOI [10.5281/zenodo.17161561](https://doi.org/10.5281/zenodo.17161561))

What You'll Learn from this Webinar

The relationship between web standards and APIs

An overview of Web API standards and their role in biomedical data services, specifically within the context of the NIAID Data Blueprint.

APIs and FAIR Principles

Leveraging standards and web architecture to ensure APIs can be expressed aligning to the FAIR principles of the NIAID Blueprint

APIs and machine readable and AI ready goals

Introduction to Model Context Protocol (MCP). Ultimately advocating for API practices that serve both human researchers and machine-driven discovery.

Audience:

Attendees range from world experts to beginners on the topic and everything in between.

Focus on key points.

Please add questions to the Q&A tool. Longer topics will shape future materials.

Play along!


Enjoy the **"code edition"** of this talk in Colab or download and run it yourself from GitHub.

[Colab Link](#) / [GitHub Link](#)

Presentation, notebook, example code and more details at:

<https://github.com/go-fair-us/apireference>
(or use QR code to the right)

Issue creation and pull requests always welcome from all!




Images throughout this presentation created with Gemini.



API Backstory: Our cast of ~~characters~~ acronyms

The *Application Programming Interface* (API)

- Started in the 60s. Described a means to interact between systems, main frames for example, and peripherals.
- This presentation, API = mostly "Web APIs"
- REST, or Representational State Transfer, from the Roy Fielding's 2000 dissertation. *It is an approach to APIs aligned with web architectural standards*




Think of Web APIs as the different ways software applications "talk" to one another. Just as we use different languages or protocols (a formal document vs. say, this presentation), developers choose different API styles based what they need.

APIs and the Blueprint

For standardized machine **access to metadata**, APIs should expose metadata elements described in Table 1. In general, an API should meet these minimum objectives:

- **Metadata Encoding:** API responses should return metadata encoded in JSON-LD... This would follow the types and properties guidance in the minimal metadata specification above.
- **IRI (URL) Structure:** API endpoints should be designed as resource-oriented IRIs (e.g., /datasets/{dataset_id}), avoiding verbs and complex query parameters in the IRI structure*. This ensures that the IRIs can function as persistent identifiers (IRIs) within the JSON-LD @id field, enabling seamless integration into knowledge graphs.
- **HTTP Method:** Metadata retrieval should be performed using the HTTP GET method.
- **Documentation:** Should adhere to OpenAPI/Swagger specifications for machine-readability and ease of use.

Though this guidance is focused on metadata, many elements apply to any resource request.




An example of API-exposed metadata formatted to meet these objectives is given in Supplemental Table 7.

<https://www.niaid.nih.gov/research/data-blueprint>

Mayer et al. (2025) DOI: [10.5281/zenodo.17161561](https://doi.org/10.5281/zenodo.17161561)





Exploring the Properties of The Request–Response

The Web, APIs and the *Request Response* Sequence

```
curl -request GET --url  
'https://import.org/data/query/api/search/study?term=influenza%20vaccine&fromRecord=0&pageSize=20&PreTag=%3Cem%3E&PostTag=%3C%2Fem%3E&format=json&sortFieldDirection=asc&conditionOrDisease=asthma%20COVID-19'
```


Same [URL](https://import.org/data/query/api/search/study?term=influenza%20vaccine&fromRecord=0&pageSize=20&PreTag=%3Cem%3E&PostTag=%3C%2Fem%3E&format=json&sortFieldDirection=asc&conditionOrDisease=asthma%20COVID-19), dropped into the browser

Interacting with APIs is typically expressed as a **requests and its response**.

Though the "curl" request looks opaque, in this case we can simply drop it into a browser and get the same response.

APIs are mostly just URLs that return formatted data as a response, vs an HTML page

```
{  
  "took": 31,  
  "timed_out": false,  
  "_shards": {  
    "total": 5,  
    "successful": 5,  
    "skipped": 0,  
    "failed": 0  
  },  
  "hits": {  
    "total": {  
      "value": 0,  
      "relation": "eq"  
    },  
    "max_score": null,  
    "hits": []  
  },  
  "aggregations": {  
    "research_focus": {  
      "doc_count_error_upper_bound": 0,  
      "sum_other_doc_count": 0,  
      "buckets": [  
        {  
          "key": "Atopy/Allergy",  
          "doc_count": 0  
        },  
        {  
          "key": "Autoimmune",  
          "doc_count": 0  
        }  
      ]  
    }  
  }  
}
```




```
{  
  "took": 29,  
  "timed_out": false,  
  "_shards": {  
    "total": 5,  
    "successful": 5,  
    "skipped": 0,  
    "failed": 0  
  },  
  "hits": {  
    "total": {  
      "value": 0,  
      "relation": "eq"  
    },  
    "max_score": null,  
    "hits": []  
  },  
  "aggregations": {  
    "research_focus": {  
      "doc_count_error_upper_bound": 0,  
      "sum_other_doc_count": 0,  
      "buckets": [  
        {  
          "key": "Atopy/Allergy",  
          "doc_count": 0  
        },  
        {  
          "key": "Autoimmune",  
          "doc_count": 0  
        },  
        {  
          "key": "Cardiovascular system disease",  
          "doc_count": 0  
        },  
        {  
          "key": "Cell Biology",  
          "doc_count": 0  
        }  
      ]  
    }  
  }  
}
```

The Web, APIs and the *Request Response* Sequence

We see this daily when we load a site
URLs into a client we call a "browser".

We use METHODS like GET, POST and
others (HEAD, PUT, DELETE, CONNECT, OPTIONS, TRACE).

Method	What it does	Real-world example
GET	Retrieves data	Viewing this presentation
POST	Sends/Creates data	Submitting a new "Sign Up" form to attend this webinar.



Example API clients include things
like curl, [HTTPIe](#), PostMan, [Bruno](#) or
custom code.

Think of these as like browsers for the
API web

Web Architecture

We can leverage the core web architecture to be an "APIs for Metadata records".


This is because URLs represent fetchable resources we can work with!

ImmPort example:

<https://www.dev.immport.org/shared/study/SDY2176/summary>

We can also fetch it via: validator.schema.org to obtain the structured JSON-LD record.

NOTE: The JSON-LD is loading dynamically into the document object model via Javascript. This means a browser or headless browser is needed to execute that process.



The screenshot displays the ImmPort website interface. At the top, there's a navigation bar with links like 'Shared Data', 'Data Catalogs', and a search bar. Below the navigation bar, a message states: 'Have questions about ImmPort? Try out ImmPort's new documentation chatbot (beta).' The main content area shows the summary for study SDY2176: 'A safe and highly efficacious measles virus-based vaccine expressing SARS-CoV-2 stabilized prefusion spike'. It includes metrics for pageviews (214) and downloads (0). A 'Download Options' section offers links for 'Study Package (Web)', 'API', and 'HL7 FHIR'. Below this, a tabbed interface shows 'Summary', 'Design', 'Adverse Event', 'Assessment', 'Interventions', 'Medications', 'Substance', 'Demographics', 'Lab Tests', 'Mechanistic Assays', and 'Study Files'. The 'Summary' tab is active, displaying a table with fields like 'Accession', 'Title', 'DOI', 'Brief Description', 'Research Focus', 'Condition Studied', 'Start Date', 'Detailed Description', 'Objectives', 'Hypothesis', 'Endpoints', 'Intervention Agent', 'Sex Included', and 'Subjects'. A red arrow points from the 'Summary' tab to the 'Dataset' tab, which is also visible. The 'Dataset' tab shows a table with columns for 'Dataset', 'Type', 'Status', 'Accession', 'Name', 'Description', and 'Accession'. It lists two datasets: 'SDY2176' and 'SDY2176-1'. The 'SDY2176' dataset is highlighted, showing its details: 'A safe and highly efficacious measles virus-based vaccine expressing SARS-CoV-2 stabilized prefusion spike'. The 'SDY2176-1' dataset is also listed, showing its details: 'A safe and highly efficacious measles virus-based vaccine expressing SARS-CoV-2 stabilized prefusion spike'. The 'Dataset' tab also displays a JSON-LD record for the study, which is a structured data format used for machine-readable metadata.

Addressing Addresses



Cool URIs

"URIs don't change: people change them."

<https://www.w3.org/Provider/Style/URI>

- From the W3C (web standards body)
- Stable, governed service URLs
 - Like GUPRIs for PIDs

Cool URIs for the Semantic Web

<https://www.w3.org/TR/cooluris/>

- Dereferenceable
 - "Be on the web", "follow your nose"
 - distinguish web docs from real-world objects
- Hash URIs (e.g., [example.com/about#alice](https://www.w3.org/TR/cooluris/))
 - Leverage hash striping, used a lot in ontologies for example
- 303 URIs (e.g., [example.com/id/alice](https://www.w3.org/TR/cooluris/)): Redirect (303 See Other) to HTML/RDF; flexible for large/dynamic data.
- Use content negotiation

Example of what a typical GET API Call resemble:



Characteristic	Cool URI Example	Not Cool URI Example
Persistence: Avoids dates (unless they have meaning) or tech specifics, or names that might change.	<code>.../people/alice</code>	<code>.../1999/alice.php</code>
Simplicity: No query strings or scripts	<code>.../books/123</code>	<code>.../cgi-bin/book.pl?id=123</code>
Meaningfulness: Readable and descriptive	<code>.../cities/paris</code>	<code>.../x7k9p2m</code>
Unicode Support (IRI): Native characters over encoding	<code>.../人名/alice</code>	<code>.../%E4%BA%BA%E5%90%8D/alice</code>

Caching

Introduction to HTTP Caching

At its core, **caching** is the practice of storing a copy of a given resource (like a JSON response or an image) and serving it back when requested again. Instead of a request traveling all the way to the application server or database, it is intercepted by a "cache" layer that stays closer to the user, providing a shortcut for data retrieval.

- Caching cuts backend compute and database load
- Caching helps services scale to more users and AI/ML workloads without proportional cost increases.
- Faster response times

NOTE: This is about server side caching, not client side. In this approach caching is an architecture component.

Better design = Better caching

- Use path for stable identifiers:
/patients/123/reports/latest
- Put filters/sorting in query only when needed (and set Cache-Control wisely)
- Clean, persistent URLs
 - a. More cache hits
 - b. Faster APIs, less load

Quick wins

- Avoid dynamic/temporary params in paths
- Use ETag/Last-Modified for conditional requests
- Set strong Cache-Control
 - a. max-age on stable resources

APIs Need Versioning

API Versioning is the practice of managing multiple versions of an API simultaneously. It allows you to introduce modern features (e.g., v2.0) while keeping the original structure (v1.0) active for those who need it.

- Versioning enables safe updates without breaking researcher tools
- Provides stable URLs for reliable long-term citations and impact
- Ensures data reproducibility across years
- Reduces maintenance effort and support tickets dramatically
- Strengthens FAIR Reusability and machine-actionability
- Future-proofs services as science and data models evolve

Types of Versioning

- URI/Path Versioning
(e.g., <https://api.data.niaid.nih.gov/v1/query>)
 - Best for Public/scientific APIs
- Header-Based Versioning (e.g., **Accept: application/vnd.myapi.v2+json** or custom header like **X-API-Version: 2**)
- Query Parameter Versioning (e.g., **/api/resources?version=2**)
 - Can make caching harder
- Media Type / Content Negotiation

There are some very robust and tested approaches to versioning like in Stripe and GitHub. An open source approach to this model can be seen in: <https://docs.cadwyn.dev/>

Formats Facilitate Correct Interpretation (Encoding)

As anywhere, incompatible formats make interoperability difficult. So standard formats returned by APIs is critical.

Examples Standards

- **Community Format Examples**
 - a. FHIR (Fast Healthcare Interoperability Resources)
 - b. GA4GH (Global Alliance for Genomics & Health)
- **CSV + Semantically Expressed Data Model:**
Even "good old" CSV is fine is out metadata can provide semantic context to the columns (see CVS on the Web, W3C)

Easier Integration & Reduced "Data Wrangling": By using common formats, scientists spend less time manually fixing errors or converting files and more time performing actual analysis.

Modern APIs use "content negotiation" to deliver the best format for the task. While a simple table (**CSV**) is popular for its ease of use, adding **JSON-LD** provides "semantic enhancement." This means the data isn't just a value; it's linked to a specific, globally defined concept.

- **The Power of PIDs (Persistent Identifiers):** PIDs ensure that data remains aligned and unambiguous over time. For example:
 - **In a standard CSV:** A column named temp could mean temperature, temporary status, or a template.
 - **In JSON-LD:** The value for temp points directly to a controlled vocabulary (e.g., "Body Temperature in Celsius"), removing all guesswork for the computer.

Metadata and Knowledge Graphs

- **JSON-LD as RDF:** This format allows data to be treated as part of a **Knowledge Graph**. Because JSON-LD is representation of RDF (Resource Description Framework), it can plug directly into evolving global networks like the **Open Knowledge Network (OKN)**, connecting disparate research findings into a single, searchable web of data.

APIs Should Use Standard Error Handling & Status Codes

Understanding HTTP Status Codes

When an API client sends a request to a server, the server doesn't just send back data; it sends a **Status Code**. These three-digit numbers act as a universal shorthand, telling the client exactly how the request was handled without needing a long text explanation.

For researchers and developers, these codes are the foundation of **Linked Data**. They provide a machine-readable way to understand if a digital object (like a dataset or a paper) has moved, is missing, or is currently unavailable.

1xx	2xx	3xx	4xx	5xx
The request was received but there isn't a response yet	The request was received successfully and the server can return the requested data	The requested resource is in another location, so a separate call is necessary	The request can't be completed due to an issue on the client's end	The request can't be completed due to an issue on the server's end

1xx → Keep going / processing: rare, 102

2xx → Success! 200 (everyone's favorite)

3xx → Go somewhere else (redirect): 301, 303

4xx → Your fault (fix the request): 404, 406

5xx → Server's fault (try later or contact support): 500, 503

Describe (Document) Your API




There are well-established approaches for describing your service in a format that's readable by both humans and machines.

One of the most popular is OpenAPI (formerly known as Swagger): <https://swagger.io/specification/>

OpenAPI = the official name of the **specification** (current versions: OpenAPI 3.0 and 3.1).

Swagger = the brand name for the popular suite of **tools** (Swagger UI, Swagger Editor, Swagger Codegen, etc.) that work with the OpenAPI Specification.



- Documentation is done in a simple text document
- Tooling provides for a common web based UI where APIs can be explored and tested directly
- Can be leveraged by machines/AI

Fun with Descriptions

Well documented APIs can be used directly with current generation AI coding agents. Leveraging good documentation to generate code and even execute results for the user.

Additionally, these can be enhanced with some skill/tool descriptions layered around them to enhance their use.

Generated Clients:


Current code generation system can produce simple clients based on OpenAPI specifications:

See: <https://github.com/go-fair-us/apireference/tree/master/code/clients/Import>

Additionally, many can execute code in an interactive "REPL" loop.

See: <https://github.com/go-fair-us/apireference/blob/master/docs/cluadeSession.md>

Example flow for AI generated clients. There would be a similar workflow for human use.



A quick Example From the Notebook

Let's revisit some of these topics we just covered with an example pulled from the notebook.

Here is a Geospatial/Geocoding example API call that scopes many of the elements we have just reviewed.

The notebook:

- 1) Makes an API call to the [OpenStreetMap Nominatim](#) (Geocoding: Address ↔ Coordinates)
- 2) Converts the results to WKT and/or GeoJSON
- 3) Plots the results with Leaflet.

This simple process leverages much of what we just reviewed and brings up some new ones.

Unofficial OpenAPI:

<https://sparkfabrik.github.io/nominatim-openapi/#/>


```
# ===== CONFIG =====
address = "Slater, Iowa"          # CHANGE THIS to anything
user_agent = "MyGeospatialAPIDemo/1.0 (drfiles@gmail.com)"
# REQUIRED by Nominatim policy – use your own name/email

# ===== API CALL SETUP =====
url = "https://nominatim.openstreetmap.org/search"

params = {
    "format": "json",              # ask for JSON
    "q": address,                  # the search query
    "limit": 1,                    # just the best match
    "addressdetails": 1            # extra info (optional but nice)
}

headers = {"User-Agent": user_agent}

response = requests.get(url, params=params, headers=headers)
```





Ecosystem of Resources

NIAID Data Ecosystem Discovery Portal Use of APIs

NDE Portal (<https://data.niaid.nih.gov/>) uses APIs and other methods for accessing metadata published by the community.

- 17 out of 41 resources catalogued in the Discovery Portal have discoverable APIs
- 37 out of 45 resources ingested by the Discovery Portal have APIs that are leveraged for the ingestion
- Primary metadata encoding format leveraged is JSON. ImmPort and others already express JSON-LD.
 - There is also XML and CSV formatted metadata records.
- Some APIs don't provide a full metadata coverage, which is then augmented with site crawling or exports from alternative API endpoints.

The screenshot shows the NIAID Data Ecosystem Discovery Portal search results page. The header includes the NIH logo and the title 'NIAID Data Ecosystem'. The search bar contains the query 'hasAPI:true'. The left sidebar shows search filters, including a date histogram for 2025 and 2026, and a list of sources. The main content area displays 'Showing results for: hasAPI:true' and 'Filtered by: Date: From 2000-01-01 to 2026-12-31'. Below this, there are three resource catalog cards: 'BacDive', 'Bacterial and Viral Bioinformatics Resource Center', and 'ClinVar'. Each card shows the date, open access status, API availability, and content types. The 'BacDive' card mentions 'The Bacterial Diversity Metadatabase BacDive is the worldwide largest database for standardized bacterial phenotypic information. Its mission is t...'. The 'Bacterial and Viral Bioinformatics Resource Center' card mentions 'The Bacterial and Viral Bioinformatics Resource Center (BV-BRC) is an information system designed to support research on bacterial and...'. The 'ClinVar' card mentions 'ClinVar is a freely accessible, public archive of reports of human variations classified for diseases and drug responses, with supporting...'. The bottom of the page shows a pagination bar with a current page indicator.

<https://data.niaid.nih.gov/search?q=hasAPI%3Atrue&filters=%28date%3A%5B%222000-01-01%22+TO+%222026-12-31%22%5D%29>

```
curl -X 'GET' \  
  'https://api.data.niaid.nih.gov/v1/query?q=hasAPI%3Atrue&facet_size=10&fetch_all=true' \  
  -H 'accept: */*'
```

Thanks to Ginger Tsueng for the material for this slide.

APIs in NIAID

When resources are online and machine readable, there are many ways it can be leveraged to provide views into the landscape of resources.

This graph is formed from the data exposed by the NDE Portal and then translated to a graph.

NOTE: The use of PIDs and vocabularies is vital in expressing networks like this.


Notice the red dot near the middle with seven connecting edges. It is:

["http://edamontology.org/topic_0622"](http://edamontology.org/topic_0622)

Which redirects (303) to
https://bioportal.bioontology.org/ontologies/EDAM?p=classes&conceptid=topic_0622

which is "Genomic"

```
curl -I http://edamontology.org/topic_0622
HTTP/1.1 303 See Other
Date: Thu, 26 Mar 2026 20:27:15 GMT
Server: Apache/2.4.63 (AlmaLinux) OpenSSL/3.5.1
X-Frame-Options: SAMEORIGIN
Location:
http://bioportal.bioontology.org/ontologies/EDAM?p=terms&conceptid=topic\_0622
Content-Type: text/html; charset=iso-8859-1
```



AI and APIs:

The Model Context Protocol (MCP)

Model Context Protocol (MCP) offers a clean, AI-ready way to expose your resources, tools, and capabilities so that AI models and agents can easily discover and use them.

It builds directly on familiar API concepts that many teams already understand, while adding the structure that AI systems need. In many ways, MCP can be thought of as **“APIs by convention for AI.”**

MCP is a stateful JSON-RPC 2.0 protocol that defines a clear, consistent set of methods. This gives AI models and agents a reliable way to:


- Discover available *tools, prompts, and resources*
- Understand exactly how to call them

Important Caution:

While MCP is powerful, it comes with unique security considerations. Teams and research groups should approach it with care and be diligent about security practices.

See more details at:

<https://github.com/go-fair-us/apireference/blob/master/code/mcp/README.md>





CAUTION: Don't just wrap your existing APIs in an MCP wrapper. You really aren't accomplishing much, other than more maintenance requirements.

MCP? Or, just APIs with descriptions or CLIs?

Do I need MCP?

Maybe, maybe not.

You might want to explore MCP if:

- If you have expressed interest from your community as an approach for access
- If you want a generic "AI Ready" service to provide a broad community via the network

If, you feel the maintenance and administration of a new service isn't justified, then you might tick the "AI Ready" box with:

- Machine readable API descriptions to allow client generation (*though not necessarily AI workflows*)
- A Command Line Interface (CLI) that an AI harness can leverage

Examples of resources that AI might be able to leverage already to address access and readiness goals.

Command Line Interface (CLI) access:

- BV-BRC: https://www.bv-brc.org/docs/cli_tutorial/index.html
- Many others via shell scripts (ie, bash + curl)
 - Recently Google released a CLI for Google Workspace (<https://github.com/googleworkspace/cli>). Why was this done? A key driver was to allow agents to leverage it.

Example OpenAPI 3.0/Swagger

- IEDB: <https://query-api.iedb.org/docs/swagger/>
- TBPortals: <https://analytic.tbportals.niaid.nih.gov/index.html>
- ImmPort: <https://immport.org/data/query/swagger-ui/index.html>
- ClinicalTrials: <https://clinicaltrials.gov/data-api/api>

APIs and FAIR Principles

The topics discussed here provide a foundation to establish a shared foundation to applying FAIR principles to APIs.

Enhanced Discoverability & Access:

- Make the API easily findable and usable by improving registration, providing clear access instructions, and ensuring persistent metadata URIs.

Machine-Readable & Interoperable Design:

- Standardize documentation (OpenAPI/Swagger)
- Through data formats (JSON-LD, CSV with Schema, FIHR, etc) to ensure seamless machine interaction and integration with other systems.
- Leverage PIDs and semantics to help provide context, to humans and machines, for these.

Consistent & Shared API Practices:

- Maintain uniformity across the API by using consistent approaches, HTTP methods, meaningful naming conventions, and aligning with established standards.

Robust Versioning & Management:

- Implement clear versioning strategies and document changes to ensure API stability and support long-term reusability.

APIs, good for humans and machines!



For the human

Things like formats, URL versioning, API approaches and explicit semantics **represent a shared context among the research community.**

For the machine

Things like formats, URL versioning, API approaches and explicit semantics **represent a shared *execution context* across distributed systems, services, and agents.**

Upcoming Webinars

Date (12 PM ET)	Topics
March 31	• Introduction to APIs Covers Web API fundamentals, including stable “Cool URIs,” interoperable formats like JSON-LD and FHIR aligned with FAIR principles, and documentation practices such as HTTP status codes and OpenAPI/Swagger. It also explores how APIs are being integrated into emerging AI and biological data ecosystems, including approaches like the Model Context Protocol (MCP).
April 29	• Improving Impact through Identifiers, Metadata and Citations Explore how persistent identifiers for research assets creates the foundation for tracking and understanding research impact. Learn strategies for developing clear, domain-aligned citation recommendations for data, and gain exposure to how identifiers are indexed along with the metrics, tools, and techniques used to measure citation impact.
May	• Increasing Findability and Interoperability: Introducing the Core Metadata Schema Introduces the core metadata schema (MS) from the Blueprint, including its key elements and the motivations for adopting it.
June	• Improving Impact through FAIR Implementation Profiles Describe and share your repo’s use of FAIR identifiers, metadata, services, and other resources.




Acknowledgements

We thank the NIH staff and community members who contributed to the Blueprint and NIAID FAIR landscaping project.

Funding Acknowledgement: This project has been funded in whole or in part with Federal funds from the National Cancer Institute, National Institutes of Health, under Contract No. 75N91019D00024, Task Order No. 75N91020F0022. The project is funded by the Frederick National Laboratory for Cancer Research, operated by Leidos Biomedical Research, Inc., on behalf of the National Institute of Allergy and Infectious Diseases.

Disclaimer: The content of this publication does not necessarily reflect the views or policies of the Department of Health and Human Services. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s).

Image source: <https://www.niaid.nih.gov/>



Discussion

APIs Need Versioning



- Versioning enables safe updates without breaking researcher tools
- Provides stable URLs for reliable long-term citations and impact
- Ensures data reproducibility across years
- Reduces maintenance effort and support tickets dramatically
- Strengthens FAIR Reusability and machine-actionability
- Future-proofs services as science and data models evolve

Types of Versioning

- URI/Path Versioning (e.g., **/api/v1/resources**, **/api/v2/resources**):
- Header-Based Versioning (e.g., **Accept: application/vnd.myapi.v2+json** or custom header like **X-API-Version: 2**)
- Query Parameter Versioning (e.g., **/api/resources?version=2**)
 - Can make caching harder
- Media Type / Content Negotiation

There are some very robust and tested approaches to versioning like in Stripe and GitHub. An open source approach to this model can be seen in: <https://docs.cadwyn.dev/>

Strategy	Pros	Cons	Best For
URI/Path	Simple, visible, cacheable	URL fragmentation	Public/scientific APIs
Header/Media Type	Clean URLs, RESTful	Less discoverable	Long-term stability (FAIR)
Query Param	Flexible	Caching issues	Quick experiments

APIs Should Use Standard Error Handling & Status Codes

Response codes communicate status between client and server

Code like 301, 404 and other are fundamental to the web and to linked data patterns.

1xx → Keep going / processing

2xx → Success!

3xx → Go somewhere else (redirect)

4xx → Your fault (fix the request)

5xx → Server's fault (try later or contact support)

1xx

The request was received but there isn't a response yet

2xx

The request was received successfully and the server can return the requested data

3xx

The requested resource is in another location, so a separate call is necessary

4xx

The request can't be completed due to an issue on the client's end

5xx

The request can't be completed due to an issue on the server's end

Status Code Examples:

200 OK: The request was successful, and the server is returning the requested resource (e.g., a web page or API data loads perfectly).

301 Moved Permanently: The requested resource has been permanently moved to a new URL, and browsers/search engines should update their links to the new location.

302 Found (or 307 Temporary Redirect): The resource is temporarily available at a different URL, so the client should use the new address for this request but keep trying the original one in the future.

404 Not Found: The server cannot find the requested resource (the classic "page not found" error you see when a link is broken or a page was deleted).

500 Internal Server Error: A generic server-side failure occurred (something broke on the website's backend, and it couldn't process your request properly).

403 Forbidden: The server understood the request but is refusing to fulfill it, usually because you lack permission to access the resource (e.g., trying to view a private admin page).

429 Too Many Requests: You (or your device/IP) have sent too many requests in a short time, hitting the server's rate limit, so it's temporarily blocking you.

Caching

Better design = Better caching

- Use path for stable identifiers: /patients/123/reports/latest
- Put filters/sorting in query only when needed (and set Cache-Control wisely)
- Clean, persistent URLs
 - a. More cache hits
 - b. Faster APIs, less load

Quick wins

- Avoid dynamic/temporary params in paths
- Use ETag/Last-Modified for conditional requests
- Set strong Cache-Control
 - a. max-age on stable resources

Some patterns might break persistence

- /report?id=42&ver=3&cachebust=abc123
 - a. Looks temporary
 - b. People hesitate to link/share

Query params

Many caches (CDNs, browsers, proxies) treat query strings as unique keys

- a. Low/no TTL or no caching at all for parameterized URLs
- b. Result: More server hits, slower responses, higher load